

January 2025
CSE 314 - OS
Sessional


Xv6 Introduction
Nafis Tahmid
May 07, 2025



Dual Mode Operation

CPU executes in 2 Modes

- **Kernel mode:** CPU can execute all machine instructions CPU can use every hardware feature
- **User mode:** permits only a subset of the instructions and a subset of the hardware features. Allows OS to protect itself and other system components.



Kernel Mode Execution

- When does CPU start executing in kernel mode?
 - A. System Boot - starting a computer
 - B. Hardware Interrupt - generated by hardware devices to signal that they need some attention from the OS.
 - C. Trap 
 - A. A software-generated interrupt caused either by an error (i.e. division by 0 or invalid memory access) or
 - B. By a specific request from a user program that an operating-system service needs to be performed.

Switching Modes

- To obtain services from the operating system,
 - an user program must make a **system call**, which **traps** into the kernel and invokes the operating system.
- The **TRAP** instruction switches from user mode to kernel mode and starts the operating system.
- When the work has been completed, control is returned to the user program at the instruction following the system call.


System Calls

- Programming interface to the services provided by the OS
- Typically used from a high-level language (C or C++) program.
- Typically a number is associated with each system call, and the OS maintains a table indexed according to these numbers.

***SYSTEM CALL**

**WHAT
USER DOES**

**WHAT IS GOING
ON INSIDE OS**



Let's see xv6

Prerequisite tools:

<https://pdos.csail.mit.edu/6.828/2022/tools.html>

Press Ctrl+A then X to quit xv6


Challenges you will face in this assignment

Too many files

Use

- `grep <pattern> *`
- `grep <pattern> kernel/*`
- `grep <pattern> user/*`

When you are the only
language people choose
to write a kernel*..



Warnings

me



xv6



✖ Cons

- Being open source, the program often crashes



Thomas had never seen such

weird
thing

before

Adding a User program

- Create a file under user folder
- These three lines must be added on top strictly following order

```
#include "kernel/types.h"  
#include "kernel/stat.h"  
#include "user/user.h"
```

Adding a User program

- Create a file under user folder
- These three lines must be added on top **strictly following order**

```
#include "kernel/types.h"  
#include "kernel/stat.h"  
#include "user/user.h"
```

User can call these functions **“declared”**
here

Adding a User program

- Create a file under user folder
- These three lines must be added on top strictly following order

```
#include "kernel/types.h"  
#include "kernel/stat.h"  
#include "user/user.h"
```

- Then add prog_name in Makefile.
- Then run make clean; make qemu

```
PROGS=\n  $U/_cat\  
  $U/_echo\  
  $U/_forktest\  
  $U/_grep\  
  $U/_init\  
  $U/_kill\  
  $U/_ln\  
  $U/_ls\  
  $U/_mkdir\  
  $U/_myprog\ # add this line  
  $U/_rm\  
  $U/_sh\  
  $U/_stressfs\  
  $U/_usertests\  
  $U/_grind\  
  $U/_wc\  
  $U/_zombie\
```

Adding a system call

First let's understand what happens during a system call.

Suppose we have added **our own** system call with signature **int getuid()**; which simply returns an **integer value** stored in a variable named **uid** in **kernel**.

gibuid.c

```
int uid = getuid();  
  
printf("%d\n", uid);
```

A series of things happen

user/gibuid.c

```
#include "kernel/types.h"
#include "kernel/stat.h"
#include "user/user.h"

int main()
{
    int uid = getuid();

    printf("%d\n", uid);

    return 0;
}
```


A series of things happen

user/gibuid.c

```
#include "kernel/types.h"
#include "kernel/stat.h"
#include "user/user.h"

int main()
{
    int uid = getuid();
    printf("%d\n", uid);

    return 0;
}
```

Signature is declared
here

user/user.h

```
user.h > ...
struct stat;

// system calls
int fork(void);
int exit(int) __attribute__((noreturn));
int wait(int*);
int pipe(int*);
int write(int, const void*, int);
int read(int, void*, int);
int close(int);
int kill(int);
int exec(const char*, char**);
int open(const char*, int);
int mknod(const char*, short, short);
int unlink(const char*);
int fstat(int fd, struct stat*);
int link(const char*, const char*);
int mkdir(const char*);
int chdir(const char*);
int dup(int);
int getpid(void);
char* sbrk(int);
int sleep(int);
int uptime(void);
int getuid(void); // signature declaration
```

A series of things happen

user/gibuid.c

What about definition?

```
#include "kernel/types.h"
#include "kernel/stat.h"
#include "user/user.h"

int main()
{
    int uid = getuid();

    printf("%d\n", uid);

    return 0;
}
```

A series of things happen

user/gibuid.c

```
#include "kernel/types.h"
#include "kernel/stat.h"
#include "user/user.h"

int main()
{
    int uid = getuid();

    printf("%d\n", uid);

    return 0;
}
```

What about definition?

Available in user/usys.S (generated from user/usys.pl)

Some other files come into play also. Let's see...


System Calls [Recap]

- Programming interface to the services provided by the OS
- Typically used from a high-level language (C or C++) program.
- Typically a number is associated with each system call, and the OS maintains a table indexed according to these numbers.

A series of things happen

user/gibuid.c

```
int uid = getuid();
```



kernel/syscall.h

```
// System call numbers
#define SYS_fork 1
#define SYS_exit 2
#define SYS_wait 3
#define SYS_pipe 4
#define SYS_read 5
#define SYS_kill 6
#define SYS_exec 7
#define SYS_fstat 8
#define SYS_chdir 9
#define SYS_dup 10
#define SYS_getpid 11
#define SYS_sbrk 12
#define SYS_sleep 13
#define SYS_uptime 14
#define SYS_open 15
#define SYS_write 16
#define SYS_mknod 17
#define SYS_unlink 18
#define SYS_link 19
#define SYS_mkdir 20
#define SYS_close 21
#define SYS_getuid 22
```

A series of things happen

user/gibuid.c

```
int uid = getuid();
```

user/usys.pl

```
# Generate usys.S, the stubs for syscalls.
```

```
4
5 print "# generated by usys.pl - do not edit\n";
```

```
6
7 print "#include \"kernel/syscall.h\"\n";
```

```
8
9 sub entry {
10     my $name = shift;
11     print ".global $name\n";
12     print "${name}:\n";
13     print "    li a7, SYS_${name}\n";
14     print "    ecall\n";
15     print "    ret\n";
16 }
```

```
34 entry("dup");
35 entry("getpid");
36 entry("sbrk");
37 entry("sleep");
38 entry("uptime");
39 entry("getuid");
40
```

kernel/syscall.h

```
// System call numbers
```

```
#define SYS_fork 1
```

```
#define SYS_exit 2
```

```
#define SYS_wait 3
```

```
#define SYS_pipe 4
```

```
#define SYS_read 5
```

```
#define SYS_kill 6
```

```
#define SYS_exec 7
```

```
#define SYS_fstat 8
```

```
#define SYS_chdir 9
```

```
#define SYS_dup 10
```

```
#define SYS_getpid 11
```

```
#define SYS_sbrk 12
```

```
#define SYS_sleep 13
```

```
#define SYS_uptime 14
```

```
#define SYS_open 15
```

```
#define SYS_write 16
```

```
#define SYS_mknod 17
```

```
#define SYS_unlink 18
```

```
#define SYS_link 19
```

```
#define SYS_mkdir 20
```

```
#define SYS_close 21
```

```
#define SYS_getuid 22
```


A series of things happen

user/gibuid.c

```
int uid = getuid();
```

user/usys.pl

```
3 # Generate usys.S, the stubs for syscalls.
4
5 print "# generated by usys.pl - do not edit\n";
6
7 print "#include \"kernel/syscall.h\"\n";
8
9 sub entry {
10     my $name = shift;
11     print ".global $name\n";
12     print "${name}:\n";
13     print "    li a7, SYS_${name}\n";
14     print "    ecall\n";
15     print "    ret\n";
16 }
17
34 entry("dup");
35 entry("getpid");
36 entry("sbrk");
37 entry("sleep");
38 entry("uptime");
39 entry("getuid");
40
```

kernel/syscall.h

```
// System call numbers
#define SYS_fork    1
#define SYS_exit    2
#define SYS_wait    3
#define SYS_pipe    4
#define SYS_read    5
#define SYS_kill    6
#define SYS_exec    7
#define SYS_fstat    8
#define SYS_chdir    9
#define SYS_dup    10
#define SYS_getpid   11
#define SYS_sbrk    12
#define SYS_sleep    13
#define SYS_uptime   14
#define SYS_open     15
#define SYS_write    16
#define SYS_mknod    17
#define SYS_unlink   18
#define SYS_link     19
#define SYS_mkdir    20
#define SYS_close    21
#define SYS_getuid   22
```

A series of things happen

user/gibuid.c

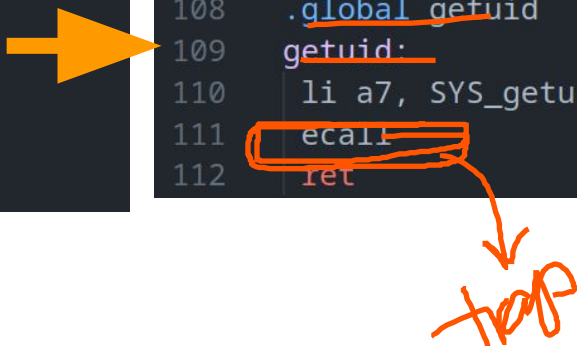
```
int uid = getuid();
```

user/usys.pl

```
3 # Generate usys.S, the stubs for syscalls.
4
5 print "# generated by usys.pl - do not edit\n";
6
7 print "#include \"kernel/syscall.h\"\n";
8
9 sub entry {
10     my $name = shift;
11     print ".global $name\n";
12     print "${name}:\n";
13     print "    li a7, SYS_${name}\n";
14     print "    ecall\n";
15     print "    ret\n";
16 }
17
34 entry("dup");
35 entry("getpid");
36 entry("sbrk");
37 entry("sleep");
38 entry("uptime");
39 entry("getuid");
40
```

user/usys.S

```
1 # generated by usys.pl - do not edit
2 #include "kernel/syscall.h"
108 .global getuid
109 getuid:
110     li a7, SYS_getuid
111     ecall
112     ret
```



kernel/syscall.h

```
// System call numbers
#define SYS_fork    1
#define SYS_exit    2
#define SYS_wait    3
#define SYS_pipe    4
#define SYS_read    5
#define SYS_kill    6
#define SYS_exec    7
#define SYS_fstat    8
#define SYS_chdir    9
#define SYS_dup    10
#define SYS_getpid   11
#define SYS_sbrk    12
#define SYS_sleep    13
#define SYS_uptime   14
#define SYS_open     15
#define SYS_write    16
#define SYS_mknod    17
#define SYS_unlink   18
#define SYS_link     19
#define SYS_mkdir    20
#define SYS_close    21
#define SYS_getuid   22
```


A series of things happen

user/gibuid.c

```
int uid = getuid();
```

Definition we are looking for..

user/usys.pl

```
3 # Generate usys.S, the stubs for syscalls.
4
5 print "# generated by usys.pl - do not edit\n";
6
7 print "#include \"kernel/syscall.h\"\n";
8
9 sub entry {
10     my $name = shift;
11     print ".global $name\n";
12     print "${name}:\n";
13     print "    li a7, SYS_${name}\n";
14     print "    ecall\n";
15     print "    ret\n";
16 }
17
34 entry("dup");
35 entry("getpid");
36 entry("sbrk");
37 entry("sleep");
38 entry("uptime");
39 entry("getuid");
40
```

user/usys.S

```
1 # generated by usys.pl - do not edit
2 #include "kernel/syscall.h"
108 .global getuid
109 getuid:
110     li a7, SYS_getuid
111     ecall
112     ret
```

Loads 22 at a7 register

kernel/syscall.h

```
// System call numbers
#define SYS_fork 1
#define SYS_exit 2
#define SYS_wait 3
#define SYS_pipe 4
#define SYS_read 5
#define SYS_kill 6
#define SYS_exec 7
#define SYS_fstat 8
#define SYS_chdir 9
#define SYS_dup 10
#define SYS_getpid 11
#define SYS_sbrk 12
#define SYS_sleep 13
#define SYS_uptime 14
#define SYS_open 15
#define SYS_write 16
#define SYS_mknod 17
#define SYS_unlink 18
#define SYS_link 19
#define SYS_mkdir 20
#define SYS_close 21
#define SYS_getuid 22
```

A series of things happen

user/gibuid.c

```
int uid = getuid();
```

Definition we are looking for..

user/usys.pl

```
3 # Generate usys.S, the stubs for syscalls.
4
5 print "# generated by usys.pl - do not edit\n";
6
7 print "#include \"kernel/syscall.h\"\n";
8
9 sub entry {
10     my $name = shift;
11     print ".global $name\n";
12     print "${name}:\n";
13     print "    li a7, SYS_${name}\n";
14     print "    ecall\n";
15     print "    ret\n";
16 }
17
34 entry("dup");
35 entry("getpid");
36 entry("sbrk");
37 entry("sleep");
38 entry("uptime");
39 entry("getuid");
40
```

user/usys.S

```
1 # generated by usys.pl - do not edit
2 #include "kernel/syscall.h"
108 .global getuid
109 getuid:
110     li a7, SYS_getuid
111     ecall
112     ret
```

Traps

Loads 22 at a7 register

kernel/syscall.h

```
// System call numbers
#define SYS_fork    1
#define SYS_exit    2
#define SYS_wait    3
#define SYS_pipe    4
#define SYS_read    5
#define SYS_kill    6
#define SYS_exec    7
#define SYS_fstat    8
#define SYS_chdir    9
#define SYS_dup    10
#define SYS_getpid   11
#define SYS_sbrk    12
#define SYS_sleep    13
#define SYS_uptime   14
#define SYS_open     15
#define SYS_write    16
#define SYS_mknod    17
#define SYS_unlink   18
#define SYS_link     19
#define SYS_mkdir    20
#define SYS_close    21
#define SYS_getuid   22
```

A series of things happen

user/gibuid.c ->

user/Usys.s ->

kernel/trampoline.S/uservec()

```
int uid = getuid();
108  .global getuid
109  getuid:
110      li a7, SYS_getuid
111      ecall
112      ret
```

```
19  trampoline:
20  .align 4
21  .globl uservec
22  uservec:
23      #
24      # trap.c sets stvec to point here, so
25      # traps from user space start here,
26      # in supervisor mode, but with a
27      # user page table.
28      #
29
30      # save user a0 in sscratch so
31      # a0 can be used to get at TRAPFRAME.
32      csrw sscratch, a0
33
34      # each process has a separate p->trapframe memory area,
35      # but it's mapped to the same virtual address
36      # (TRAPFRAME) in every process's user page table.
37      li a0, TRAPFRAME
38
39      # save the user registers in TRAPFRAME
40      sd ra, 40(a0)
41      sd sp, 48(a0)
42      sd gp, 56(a0)
43      sd tp, 64(a0)
44      sd t0, 72(a0)
45      sd t1, 80(a0)
46      sd t2, 88(a0)
47      sd s0, 96(a0)
48      sd s1, 104(a0)
49      sd a1, 120(a0)
50      sd a2, 128(a0)
51      sd a3, 136(a0)
52      sd a4, 144(a0)
53      sd a5, 152(a0)
54      sd a6, 160(a0)
55      sd a7, 168(a0)
56      sd s2, 176(a0)
57      sd s3, 184(a0)
```

A series of things happen

user/gibuid.c ->

user/Usys.s ->

kernel/trampoline.S/uservec()

```
int uid = getuid();
108     .global getuid
109     getuid:
110     li a7, SYS_getuid
111     ecall
112     ret
```

```
19     trampoline:
20     .align 4
21     .globl uservec
22     uservec:
23     #
24     # trap.c sets stvec to point here, so
25     # traps from user space start here,
26     # in supervisor mode, but with a
27     # user page table.
28     #
29
30     # save user a0 in sscratch so
31     # a0 can be used to get at TRAPFRAME.
32     csrr sscratch, a0
33
34     # each process has a separate p->trapframe memory area,
35     # but it's mapped to the same virtual address
36     # (TRAPFRAME) in every process's user page table.
37     li a0, TRAPFRAME
38
39     # save the user registers in TRAPFRAME
40     sd ra, 40(a0)
41     sd sp, 48(a0)
42     sd gp, 56(a0)
43
44     # jump to usertrap(), which does not return
45     jr t0
46
47     su s0, 20(a0)
48     sd s1, 104(a0)
49     sd a1, 120(a0)
50     sd a2, 128(a0)
51     sd a3, 136(a0)
52     sd a4, 144(a0)
53     sd a5, 152(a0)
54     sd a6, 160(a0)
55     sd a7, 168(a0)
56     sd s2, 176(a0)
57     sd s3, 184(a0)
```

A series of things happen

user/gibuid.c -> user/Usys.s -> kernel/trampoline.S/uservec() -> kernel/trap.c/usertrap()

```
int uid = getuid(); 108 .global getuid
109 getuid:
110 li a7, SYS_getuid
111 ecall
112 ret
```

```
33 // handle an interrupt, exception, or system call from user space.
34 // called from trampoline.S
35 //
36 void
37 usertrap(void)
38 {
39     int which_dev = 0;
40
41     if((r_sstatus() & SSTATUS_SPP) != 0)
42         panic("usertrap: not from user mode");
43
44     // send interrupts and exceptions to kerneltrap(),
45     // since we're now in the kernel.
46     w_stvec((uint64)kernelvec);
47
48     struct proc *p = myproc();
49
50     // save user program counter.
51     p->trapframe->epc = r_sepc();
52
53     if(r_scause() == 8){
54         // system call
55
56         if(killed(p))
57             exit(-1);
58
59         // sepc points to the ecall instruction,
60         // but we want to return to the next instruction.
61         p->trapframe->epc += 4;
62
63         // an interrupt will change sepc, scause, and sstatus,
64         // so enable only now that we're done with those registers.
65         intr_on();
66
67         syscall();
```

```
19 trampoline:
20 .align 4
21 .globl uservec
22 uservec:
23 #
24 # trap.c sets stvec to point here, so
25 # traps from user space start here,
26 # in supervisor mode, but with a
27 # user page table.
28 #
29
30 # save user a0 in sscratch so
31 # a0 can be used to get at TRAPFRAME.
32 csrw sscratch, a0
33
34 # each process has a separate p->trapframe memory area,
35 # but it's mapped to the same virtual address
36 # (TRAPFRAME) in every process's user page table.
37 li a0, TRAPFRAME
38
39 # save the user registers in TRAPFRAME
40 sd ra, 40(a0)
41 sd sp, 48(a0)
42 sd gp, 56(a0)
43
44 # jump to usertrap(), which does not return
45 jr t0
46
47
48 sd s0, 20(a0)
49 sd s1, 104(a0)
50 sd a1, 120(a0)
51 sd a2, 128(a0)
52 sd a3, 136(a0)
53 sd a4, 144(a0)
54 sd a5, 152(a0)
55 sd a6, 160(a0)
56 sd a7, 168(a0)
57 sd s2, 176(a0)
58 sd s3, 184(a0)
```


A series of things happen

user/gibuid.c → user/Usys.s → kernel/trampoline.S/uservec() → kernel/trap.c/usertrap() → syscall.c/syscall()

```
int uid = getuid(); 108 .global getuid
109 getuid:
110 li a7, SYS_getuid
111 ecall
112 ret
```

```
33 // handle an interrupt, exception, or system call from user space.
34 // called from trampoline.S
35 //
36 void
37 usertrap(void)
38 {
39     int which_dev = 0;
40
41     if((r_sstatus() & SSTATUS_SPP) != 0)
42         panic("usertrap: not from user mode");
43
44     // send interrupts and exceptions to kerneltrap(),
45     // since we're now in the kernel.
46     w_stvec((uint64)kernelvec);
47
48     struct proc *p = myproc();
49
50     // save user program counter.
51     p->trapframe->epc = r_sepc();
52
53     if(r_scause() == 8){
54         // system call
55
56         if(killed(p))
57             exit(-1);
58
59         // sepc points to the ecall instruction,
60         // but we want to return to the next instruction.
61         p->trapframe->epc += 4;
62
63         // an interrupt will change sepc, scause, and sstatus,
64         // so enable only now that we're done with those registers.
65         intr_on();
66
67         syscall();
```

kernel/syscall.c

```
135 void
136 syscall(void)
137 {
138     int num;
139     struct proc *p = myproc();
140
141     num = p->trapframe->a7;
142     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143         // Use num to lookup the system call function for num, call it,
144         // and store its return value in p->trapframe->a0
145         p->trapframe->a0 = syscalls[num]();
146     } else {
147         printf("%d %s: unknown sys call %d\n",
148             p->pid, p->name, num);
149         p->trapframe->a0 = -1;
150     }
151 }
```

A series of things happen

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->syscall.c/syscall()

```
int uid = getuid(); 108 .global getuid
                    109 getuid:
                    110 li a7, SYS_getuid
                    111 ecall
                    112 ret
```

kernel/proc.h

```
82 enum procstate { UNUSED, USED, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
83
84 // Per-process state
85 struct proc {
86     struct spinlock lock;
87
88     // p->lock must be held when using these:
89     enum procstate state; // Process state
90     void *chan;           // If non-zero, sleeping on chan
91     int killed;           // If non-zero, have been killed
92     int xstate;           // Exit status to be returned to parent's wait
93     int pid;              // Process ID
94
95     // wait_lock must be held when using this:
96     struct proc *parent; // Parent process
97
98     // these are private to the process, so p->lock need not be held.
99     uint64 kstack;        // Virtual address of kernel stack
100     uint64 sz;            // Size of process memory (bytes)
101     pagetable_t pagetable; // User page table
102     struct trapframe *trapframe; // data page for trampoline.S
103     struct context context; // switch() here to run process
104     struct file *ofile[NOFILE]; // Open files
105     struct inode *cwd;        // Current directory
106     char name[16];           // Process name (debugging)
107 };
108
```

kernel/syscall.c

```
135 void
136 syscall(void)
137 {
138     int num;
139     struct proc *p = myproc();
140
141     num = p->trapframe->a7;
142     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143         // Use num to lookup the system call function for num, call it,
144         // and store its return value in p->trapframe->a0
145         p->trapframe->a0 = syscalls[num]();
146     } else {
147         printf("%d %s: unknown sys call %d\n",
148             p->pid, p->name, num);
149         p->trapframe->a0 = -1;
150     }
151 }
```

A series of things happen

user/gibuid.c->user/Usys.s ->**kernel/trampoline.S/uservec()**->kernel/t

```
int uid = getuid(); 108 .global getuid
109 getuid:
110 li a7, SYS_getuid
111 ecall
112 ret
```

kernel/proc.h

```
82 enum procstate { UNUSED, USED, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
83
84 // Per-process state
85 struct proc {
86     struct spinlock lock;
87
88     // p->lock must be held when using these:
89     enum procstate state; // Process state
90     void *chan; // If non-zero, sleeping on chan
91     int killed; // If non-zero, have been killed
92     int xstate; // Exit status to be returned to parent's wait
93     int pid; // Process ID
94
95     // wait_lock must be held when using this:
96     struct proc *parent; // Parent process
97
98     // these are private to the process, so p->lock need not be held.
99     uint64 kstack; // Virtual address of kernel stack
100     uint64 sz; // Size of process memory (bytes)
101     pagetable_t pagetable; // User page table
102     struct trapframe *trapframe; // data page for trampoline.S
103     struct context context; // switch() here to run process
104     struct file *ofile[NOFILE]; // Open files
105     struct inode *cwd; // Current directory
106     char name[16]; // Process name (debugging)
107 };
108
```

```
43 struct trapframe {
44     /* 0 */ uint64 kernel_satp; // kernel page table
45     /* 8 */ uint64 kernel_sp; // top of process's kernel stack
46     /* 16 */ uint64 kernel_trap; // usertrap()
47     /* 24 */ uint64 epc; // saved user program counter
48     /* 32 */ uint64 kernel_hartid; // saved kernel tp
49     /* 40 */ uint64 ra;
50     /* 48 */ uint64 sp;
51     /* 56 */ uint64 gp;
52     /* 64 */ uint64 tp;
53     /* 72 */ uint64 t0;
54     /* 80 */ uint64 t1;
55     /* 88 */ uint64 t2;
56     /* 96 */ uint64 s0;
57     /* 104 */ uint64 s1;
58     /* 112 */ uint64 a0;
59     /* 120 */ uint64 a1;
60     /* 128 */ uint64 a2;
61     /* 136 */ uint64 a3;
62     /* 144 */ uint64 a4;
63     /* 152 */ uint64 a5;
64     /* 160 */ uint64 a6;
65     /* 168 */ uint64 a7;
66     /* 176 */ uint64 s2;
67     /* 184 */ uint64 s3;
68     /* 192 */ uint64 s4;
69     /* 200 */ uint64 s5;
70     /* 208 */ uint64 s6;
71     /* 216 */ uint64 s7;
72     /* 224 */ uint64 s8;
73     /* 232 */ uint64 s9;
74     /* 240 */ uint64 s10;
75     /* 248 */ uint64 s11;
76     /* 256 */ uint64 t3;
77     /* 264 */ uint64 t4;
78     /* 272 */ uint64 t5;
79     /* 280 */ uint64 t6;
80 };
```

Trap Frames are used to store the **registers** of the **current thread** when an **interrupt** occurs, or when there is a **system service call** (the **transfer** from user mode to kernel mode). Loaded in **trampoline.S**

A series of things happen

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->syscall.c/syscall()

```
int uid = getuid(); 108 .global getuid
                   109 getuid:
                   110 li a7, SYS_getuid
                   111 ecall
                   112 ret
```

kernel/proc.h

```
82 enum procstate { UNUSED, USED, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
83
84 // Per-process state
85 struct proc {
86     struct spinlock lock;
87
88     // p->lock must be held when using these:
89     enum procstate state; // Process state
90     void *chan;           // If non-zero, sleeping on chan
91     int killed;           // If non-zero, have been killed
92     int xstate;           // Exit status to be returned to parent's wait
93     int pid;              // Process ID
94
95     // wait_lock must be held when using this:
96     struct proc *parent; // Parent process
97
98     // these are private to the process, so p->lock need not be held.
99     uint64 kstack;        // Virtual address of kernel stack
100    uint64 sz;             // Size of process memory (bytes)
101    pagetable_t pagetable; // User page table
102    struct trapframe *trapframe; // data page for trampoline.S
103    struct context context; // switch() here to run process
104    struct file *ofile[NOFILE]; // Open files
105    struct inode *cwd;        // Current directory
106    char name[16];           // Process name (debugging)
107 };
108
```

Process Control Block (PCB)

Process management	Memory management	File management
Registers	Pointer to text segment	Root directory
Program counter	Pointer to data segment	Working directory
Program status word	Pointer to stack segment	File descriptors
Stack pointer		User ID
Process state		Group ID
Priority		
Scheduling parameters		
Process ID		
Parent process		
Process group		
Signals		
Time when process started		
CPU time used		
Children's CPU time		
Time of next alarm		

Figure: Fields of a PCB

A series of things happen


user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->syscall.c/syscall()

```
int uid = getuid(); 108 .global getuid
109 getuid:
110 li a7, SYS_getuid
111 ecall
112 ret
```

kernel/proc.h

```
82 enum procstate { UNUSED, USED, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
83
84 // Per-process state
85 struct proc {
86     struct spinlock lock;
87
88     // p->lock must be held when using these:
89     enum procstate state; // Process state
90     void *chan; // If non-zero, sleeping on chan
91     int killed; // If non-zero, have been killed
92     int xstate; // Exit status to be returned to parent's wait
93     int pid; // Process ID
94
95     // wait_lock must be held when using this:
96     struct proc *parent; // Parent process
97
98     // these are private to the process, so p->lock need not be held.
99     uint64 kstack; // Virtual address of kernel stack
100     uint64 sz; // Size of process memory (bytes)
101     pagetable_t pagetable; // User page table
102     struct trapframe *trapframe; // data page for trampoline.S
103     struct context context; // switch() here to run process
104     struct file *ofile[NOFILE]; // Open files
105     struct inode *cwd; // Current directory
106     char name[16]; // Process name (debugging)
107 };
108
```

Multithreaded Processes



A series of things happen

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->syscall.c/syscall()

```
int uid = getuid(); 108 .global getuid
109 getuid:
110 li a7, SYS_getuid
111 ecall
112 ret
```

kernel/syscall.c

```
107 // An array mapping syscall numbers from syscall.h
108 // to the function that handles the system call.
109 static uint64 (*syscalls[])(void) = {
110 [SYS_fork] sys_fork,
111 [SYS_exit] sys_exit,
112 [SYS_wait] sys_wait,
113 [SYS_pipe] sys_pipe,
114 [SYS_read] sys_read,
115 [SYS_kill] sys_kill,
116 [SYS_exec] sys_exec,
117 [SYS_fstat] sys_fstat,
118 [SYS_chdir] sys_chdir,
119 [SYS_dup] sys_dup,
120 [SYS_getpid] sys_getpid,
121 [SYS_sbrk] sys_sbrk,
122 [SYS_sleep] sys_sleep,
123 [SYS_uptime] sys_uptime,
124 [SYS_open] sys_open,
125 [SYS_write] sys_write,
126 [SYS_mknod] sys_mknod,
127 [SYS_unlink] sys_unlink,
128 [SYS_link] sys_link,
129 [SYS_mkdir] sys_mkdir,
130 [SYS_close] sys_close,
131 [SYS_getuid] sys_getuid,
132 };
```

kernel/syscall.c

```
135 void
136 syscall(void)
137 {
138     int num;
139     struct proc *p = myproc();
140
141     num = p->trapframe->a7;
142     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143         // Use num to lookup the system call function for num, call it,
144         // and store its return value in p->trapframe->a0
145         p->trapframe->a0 = syscalls[num]();
146     } else {
147         printf("%d %s: unknown sys call %d\n",
148             p->pid, p->name, num);
149         p->trapframe->a0 = -1;
150     }
151 }
```

A series of things happen

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->syscall.c/syscall()

```
int uid = getuid(); 108 .global getuid
109 getuid:
110 li a7, SYS_getuid
111 ecall
112 ret
```

kernel/syscall.c

```
107 // An array mapping syscall numbers from syscall.h
108 // to the function that handles the system call.
109 static uint64 (*syscalls[])(void) = {
110 [SYS_fork] sys_fork,
111 [SYS_exit] sys_exit,
112 [SYS_wait] sys_wait,
113 [SYS_pipe] sys_pipe,
114 [SYS_read] sys_read,
115 [SYS_kill] sys_kill,
116 [SYS_exec] sys_exec,
117 [SYS_fstat] sys_fstat,
118 [SYS_chdir] sys_chdir,
119 [SYS_dup] sys_dup,
120 [SYS_getpid] sys_getpid,
121 [SYS_sbrk] sys_sbrk,
122 [SYS_sleep] sys_sleep,
123 [SYS_uptime] sys_uptime,
124 [SYS_open] sys_open,
125 [SYS_write] sys_write,
126 [SYS_mknod] sys_mknod,
127 [SYS_unlink] sys_unlink,
128 [SYS_link] sys_link,
129 [SYS_mkdir] sys_mkdir,
130 [SYS_close] sys_close,
131 [SYS_getuid] sys_getuid,
132 };
133
```

kernel/syscall.c

Defined mostly in sysproc.c and sysfile.c

kernel/syscall.c

```
135 void
136 syscall(void)
137 {
138 int num;
139 struct proc *p = myproc();
140
141 num = p->trapframe->a7;
142 if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143 // Use num to lookup the system call function for num, call it,
144 // and store its return value in p->trapframe->a0
145 p->trapframe->a0 = syscalls[num]();
146 } else {
147 printf("%d %s: unknown sys call %d\n",
148 p->pid, p->name, num);
149 p->trapframe->a0 = -1;
150 }
151 }
```

kernel/defs.h

```
190 // number of elements in fixed-size array
191 #define NELEM(x) (sizeof(x)/sizeof((x)[0]))
192
```


A series of things happen

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->syscall.c/syscall()

```
int uid = getuid(); 108 .global getuid
109 getuid:
110 li a7, SYS_getuid
111 ecall
112 ret
```

kernel/syscall.c

```
107 // An array mapping syscall numbers from syscall.h
108 // to the function that handles the system call.
109 static uint64 (*syscalls[])(void) = {
110 [SYS_fork] sys_fork,
111 [SYS_exit] sys_exit,
112 [SYS_wait] sys_wait,
113 [SYS_pipe] sys_pipe,
114 [SYS_read] sys_read,
115 [SYS_kill] sys_kill,
116 [SYS_exec] sys_exec,
117 [SYS_fstat] sys_fstat,
118 [SYS_chdir] sys_chdir,
119 [SYS_dup] sys_dup,
120 [SYS_getpid] sys_getpid,
121 [SYS_sbrk] sys_sbrk,
122 [SYS_sleep] sys_sleep,
123 [SYS_uptime] sys_uptime,
124 [SYS_open] sys_open,
125 [SYS_write] sys_write,
126 [SYS_mknod] sys_mknod,
127 [SYS_unlink] sys_unlink,
128 [SYS_link] sys_link,
129 [SYS_mkdir] sys_mkdir,
130 [SYS_close] sys_close,
131 [SYS_getuid] sys_getuid,
132 };
133
134 // Prototypes for the functions that handle system calls.
135 extern uint64 sys_fork(void);
136 extern uint64 sys_exit(void);
137 extern uint64 sys_wait(void);
138 extern uint64 sys_pipe(void);
139 extern uint64 sys_read(void);
140 extern uint64 sys_kill(void);
141 extern uint64 sys_exec(void);
142 extern uint64 sys_fstat(void);
143 extern uint64 sys_chdir(void);
144 extern uint64 sys_dup(void);
145 extern uint64 sys_getpid(void);
146 extern uint64 sys_sbrk(void);
147 extern uint64 sys_sleep(void);
148 extern uint64 sys_uptime(void);
149 extern uint64 sys_open(void);
150 extern uint64 sys_write(void);
151 extern uint64 sys_mknod(void);
152 extern uint64 sys_unlink(void);
153 extern uint64 sys_link(void);
154 extern uint64 sys_mkdir(void);
155 extern uint64 sys_close(void);
156 extern uint64 sys_getuid(void);
```

kernel/syscall.c

Defined mostly in sysproc.c and sysfile.c

These were defined in syscall.h (loaded in a7). So, line 145 will call sys_getuid() [defined in sysproc.c]

kernel/syscall.c

```
135 void
136 syscall(void)
137 {
138     int num;
139     struct proc *p = myproc();
140
141     num = p->trapframe->a7;
142     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143         // Use num to lookup the system call function for num, call it,
144         // and store its return value in p->trapframe->a0
145         p->trapframe->a0 = syscalls[num]();
146     } else {
147         printf("%d %s: unknown sys call %d\n",
148             p->pid, p->name, num);
149         p->trapframe->a0 = -1;
150     }
151 }
```

kernel/defs.h

```
190 // number of elements in fixed-size array
191 #define NELEM(x) (sizeof(x)/sizeof((x)[0]))
192
```

A series of things happen

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->syscall.c/syscall()->kernel/sysproc.c

```
int uid = getuid(); 108 .global getuid
109 getuid:
110 li a7, SYS_getuid
111 ecall
112 ret
```

kernel/sysproc.c

```
95 uint64
96 sys_getuid(void)
97 {
98     return getuid();
99 }
```

These were defined in syscall.h (loaded in a7). So, line 145 will call sys_getuid() [defined in sysproc.c]

kernel/syscall.c

```
135 void
136 syscall(void)
137 {
138     int num;
139     struct proc *p = myproc();
140
141     num = p->trapframe->a7;
142     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143         // Use num to lookup the system call function for num, call it,
144         // and store its return value in p->trapframe->a0
145         p->trapframe->a0 = syscalls[num]();
146     } else {
147         printf("%d %s: unknown sys call %d\n",
148             p->pid, p->name, num);
149         p->trapframe->a0 = -1;
150     }
151 }
```

kernel/defs.h

```
190 // number of elements in fixed-size array
191 #define NELEM(x) (sizeof(x)/sizeof((x)[0]))
192
```

A series of things happen

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->syscall.c/syscall()->kernel/sysproc.c
->kernel/proc.c

```
108  .global getuid
109  getuid:
110  li a7, SYS_getuid
111  ecall
112  ret
```

kernel/sysproc.c

```
95  uint64
96  sys_getuid(void)
97  {
98  |   return getuid();
99  }
```

kernel/proc.c

```
697  int uid=123;
698
699  int getuid(void) {
700  |   return uid;
701  }
```



kernel/syscall.c

```
135  void
136  syscall(void)
137  {
138  |   int num;
139  |   struct proc *p = myproc();
140
141  |   num = p->trapframe->a7;
142  |   if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143  |       // Use num to lookup the system call function for num, call it,
144  |       // and store its return value in p->trapframe->a0
145  |       p->trapframe->a0 = syscalls[num]();
146  |   } else {
147  |       printf("%d %s: unknown sys call %d\n",
148  |           p->pid, p->name, num);
149  |       p->trapframe->a0 = -1;
150  |   }
151  }
```

kernel/defs.h

```
190  // number of elements in fixed-size array
191  #define NELEM(x) (sizeof(x)/sizeof((x)[0]))
192
```


A series of things happen

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->syscall.c/syscall()->kernel/sysproc.c
->kernel/proc.c

```
108  .global getuid
109  getuid:
110      li a7, SYS_getuid
111      ecall
112      ret
```

kernel/sysproc.c

```
95  uint64
96  sys_getuid(void)
97  {
98      return getuid();
99  }
```

kernel/proc.c

```
697  int uid=123;
698
699  int getuid(void) {
700      return uid;
701  }
```

to make getuid() accessible from kernel/sysproc.c, add this **declaration** in **kernel/defs.h**. If no new function was defined, this line wouldn't have been required.

```
107  int  ether_copyin(void *, void *, int);
108  void  procdump(void);
109  int  getuid(void);
```

kernel/syscall.c

```
135  void
136  syscall(void)
137  {
138      int num;
139      struct proc *p = myproc();
140
141      num = p->trapframe->a7;
142      if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143          // Use num to lookup the system call function for num, call it,
144          // and store its return value in p->trapframe->a0
145          p->trapframe->a0 = syscalls[num]();
146      } else {
147          printf("%d %s: unknown sys call %d\n",
148              p->pid, p->name, num);
149          p->trapframe->a0 = -1;
150      }
151  }
```

kernel/defs.h

```
190  // number of elements in fixed-size array
191  #define NELEM(x) (sizeof(x)/sizeof((x)[0]))
192
```

A series of things happen

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->syscall.c/syscall()->kernel/sysproc.c

->kernel/proc.c

kernel/syscall.c

```
108  .global getuid
109  getuid:
110      li a7, SYS_getuid
111      ecall
112      ret
```

kernel/sysproc.c

```
95  uint64
96  sys_getuid(void)
97  {
98      return getuid();
99  }
```

kernel/proc.c

```
697  int uid=123;
698
699  int getuid(void) {
700      return uid;
701  }
```

Why bother adding another function in proc.c?

Ans: It's the style of the coding in xv6, treating sysproc.c as a mediator having sys call handlers who verifies input and then delegates. We should adhere to the practice.

```
135  void
136  syscall(void)
137  {
138      int num;
139      struct proc *p = myproc();
140
141      num = p->trapframe->a7;
142      if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143          // Use num to lookup the system call function for num, call it,
144          // and store its return value in p->trapframe->a0
145          p->trapframe->a0 = syscalls[num]();
146      } else {
147          printf("%d %s: unknown sys call %d\n",
148              p->pid, p->name, num);
149          p->trapframe->a0 = -1;
150      }
151  }
```

kernel/defs.h

```
107  int  ether_copyin(void *, void *, int);
108  void  procdump(void);
109  int  getuid(void);
```

```
190  // number of elements in fixed-size array
191  #define NELEM(x) (sizeof(x)/sizeof((x)[0]))
192
```

A series of things happened(now returning)

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->syscall.c/syscall()->kernel/sysproc.c

```
int uid = getuid(); 108  .global getuid
                    109  getuid:
                    110  li a7, SYS_getuid
                    111  ecall
                    112  ret
```

kernel/sysproc.c

```
95  uint64
96  sys_getuid(void)
97  {
98      return getuid();
99  }
```

kernel/syscall.c

```
135 void
136 syscall(void)
137 {
138     int num;
139     struct proc *p = myproc();
140
141     num = p->trapframe->a7;
142     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143         // Use num to lookup the system call function for num, call it,
144         // and store its return value in p->trapframe->a0
145         p->trapframe->a0 = syscalls[num]();
146     } else {
147         printf("%d %s: unknown sys call %d\n",
148             p->pid, p->name, num);
149         p->trapframe->a0 = -1;
150     }
151 }
```

kernel/defs.h

```
190 // number of elements in fixed-size array
191 #define NELEM(x) (sizeof(x)/sizeof((x)[0]))
192
```

A series of things happened(now returning)

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->syscall.c/syscall()

```
int uid = getuid(); 108  .global getuid
                    109  getuid:
                    110  li a7, SYS_getuid
                    111  ecall
                    112  ret
```

kernel/syscall.c

```
135 void
136 syscall(void)
137 {
138     int num;
139     struct proc *p = myproc();
140
141     num = p->trapframe->a7;
142     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143         // Use num to lookup the system call function for num, call it,
144         // and store its return value in p->trapframe->a0
145         p->trapframe->a0 = syscalls[num]();
146     } else {
147         printf("%d %s: unknown sys call %d\n",
148             p->pid, p->name, num);
149         p->trapframe->a0 = -1;
150     }
151 }
```

kernel/defs.h

```
190 // number of elements in fixed-size array
191 #define NELEM(x) (sizeof(x)/sizeof((x)[0]))
192
```

A series of things happenIng

user/gibuid.c->user/Usys.s->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()

```
int uid = getuid(); 108 .global getuid
109 getuid:
110 li a7, SYS_getuid
111 ecall
112 ret
```

```
36 void
37 usertrap(void)
38 {
39     int which_dev = 0;
40
41     if((r_sstatus() & SSTATUS_SPP) != 0)
42         panic("usertrap: not from user mode");
43
44     // send interrupts and exceptions to kerneltrap(),
45     // since we're now in the kernel.
46     w_stvec((uint64)kernelvec);
47
48     struct proc *p = myproc();
49
50     // save user program counter.
51     p->trapframe->epc = r_sepc();
52
53     if(r_scause() == 8){
54         // system call
55
56         if(killed(p))
57             exit(-1);
58
59         // sepc points to the ecall instruction,
60         // but we want to return to the next instruction.
61         p->trapframe->epc += 4;
62
63         // an interrupt will change sepc, scause, and sstat
64         // so enable only now that we're done with those registers
65         intr_on();
66         syscall();
67     } else if((which_dev = devintr()) != 0){
68         // ok
69     } else {
70         printf("usertrap(): unexpected scause 0x%x pid=%d\n", r_scause(), p->pid);
71         printf("             sepc=0x%x stval=0x%x\n", r_sepc(), r_stval);
72         setkilled(p);
73     }
74
75     if(killed(p))
76         exit(-1);
77
78     // give up the CPU if this is a timer interrupt.
79     if(which_dev == 2)
80         yield();
81
82     usertrapret();
83 }
84
85
```

A series of things happenIng

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->

kernel/trap.c/usertrapret()

```
int uid = getuid();
```

```
89 void
90 usertrapret(void)
91 {
92     struct proc *p = myproc();
93
94     // we're about to switch the destination of traps from
95     // kerneltrap() to usertrap(), so turn off interrupts until
96     // we're back in user space, where usertrap() is correct.
97     intr_off();
98
99     // send syscalls, interrupts, and exceptions to uservec in tramp
100     uint64 trampoline_uservec = TRAMPOLINE + (uservec - trampoline);
101     w_stvec(trampoline_uservec);
102
103     // set up trapframe values that uservec will need when
104     // the process next traps into the kernel.
105     p->trapframe->kernel_satp = r_satp(); // kernel page table
106     p->trapframe->kernel_sp = p->kstack + PGSIZE; // process's kernel
107     p->trapframe->kernel_trap = (uint64)usertrap;
108     p->trapframe->kernel_hartid = r_tp(); // hartid for cpuid
109
110     // set up the registers that trampoline.S's sret will use
111     // to get to user space.
112
113     // set S Previous Privilege mode to User.
114     unsigned long x = r_sstatus();
115     x &= ~SSTATUS_SPP; // clear SPP to 0 for user mode
116     x |= SSTATUS_SPIE; // enable interrupts in user mode
117     w_sstatus(x);
118
119     // set S Exception Program Counter to the saved user pc.
120     w_sepc(p->trapframe->epc);
121
122     // tell trampoline.S the user page table to switch to.
123     uint64 satp = MAKE_SATP(p->pagetable);
124
125     // jump to userret in trampoline.S at the top of memory, which
126     // switches to the user page table, restores user registers,
127     // and switches to user mode with sret.
128     uint64 trampoline_userret = TRAMPOLINE + (userret - trampoline);
129     ((void (*)(uint64))trampoline_userret)(satp);
130 }
131
```

```
36 void
37 usertrap(void)
38 {
39     int which_dev = 0;
40
41     if((r_sstatus() & SSTATUS_SPP) != 0)
42         panic("usertrap: not from user mode");
43
44     // send interrupts and exceptions to kerneltrap(),
45     // since we're now in the kernel.
46     w_stvec((uint64)kernelvec);
47
48     struct proc *p = myproc();
49
50     // save user program counter.
51     p->trapframe->epc = r_sepc();
52
53     if(r_scause() == 8){
54         // system call
55
56         if(killed(p))
57             exit(-1);
58
59         // sepc points to the ecall instruction,
60         // but we want to return to the next instruction.
61         p->trapframe->epc += 4;
62
63         // an interrupt will change sepc, scause, and sstat
64         // so enable only now that we're done with those re
65         intr_on();
66
67         syscall();
68     } else if((which_dev = devintr()) != 0){
69         // ok
70     } else {
71         printf("usertrap(): unexpected scause 0x%x pid=%d\n",
72                r_scause(), p->pid);
73         printf("             sepc=0x%x stval=0x%x\n", r_sepc(),
74                r_stval());
75         setkilled(p);
76     }
77
78     if(killed(p))
79         exit(-1);
80
81     // give up the CPU if this is a timer interrupt.
82     if(which_dev == 2)
83         yield();
84
85     usertrapret();
86 }
87
```


A series of things happenIng

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->kernel/trampoline.S/userret()

```
int uid = getuid();
```

```
89 void
90 usertrapret(void)
91 {
92     struct proc *p = myproc();
93
94     // we're about to switch the destination of traps from
95     // kerneltrap() to usertrap(), so turn off interrupts until
96     // we're back in user space, where usertrap() is correct.
97     intr_off();
98
99     // send syscalls, interrupts, and exceptions to uservec in trampoline.S
100     uint64 trampoline_uservec = TRAMPOLINE + (uservec - trampoline);
101     w_stvec(trampoline_uservec);
102
103     // set up trapframe values that uservec will need when
104     // the process next traps into the kernel.
105     p->trapframe->kernel_satp = r_satp(); // kernel page table
106     p->trapframe->kernel_sp = p->kstack + PGSIZE; // process's kernel stack
107     p->trapframe->kernel_trap = (uint64)usertrap;
108     p->trapframe->kernel_hartid = r_tp(); // hartid for cpuid
109
110     // set up the registers that trampoline.S's sret will use
111     // to get to user space.
112
113     // set S Previous Privilege mode to User.
114     unsigned long x = r_sstatus();
115     x &= ~SSTATUS_SPP; // clear SPP to 0 for user mode
116     x |= SSTATUS_SPIE; // enable interrupts in user mode
117     w_sstatus(x);
118
119     // set S Exception Program Counter to the saved user pc.
120     w_sepc(p->trapframe->epc);
121
122     // tell trampoline.S the user page table to switch to.
123     uint64 satp = MAKE_SATP(p->pagetable);
124
125     // jump to userret in trampoline.S at the top of memory, which
126     // switches to the user page table, restores user registers,
127     // and switches to user mode with sret.
128     uint64 trampoline_userret = TRAMPOLINE + (userret - trampoline);
129     ((void (*)(uint64))trampoline_userret)(satp);
130 }
131
```

kernel/trampoline.S

```
100 .globl userret
101 userret:
102     # userret(pagetable)
103     # called by usertrapret() in trap.c to
104     # switch from kernel to user.
105     # a0: user page table, for satp.
106
107     # switch to the user page table.
108     sfence.vma zero, zero
109     csrw satp, a0
110     sfence.vma zero, zero
111
112     li a0, TRAPFRAME
113
114     # restore all but a0 from TRAPFRAME
115     ld ra, 40(a0)
116     ld sp, 48(a0)
117     ld gp, 56(a0)
118     ld tp, 64(a0)
119     ld t0, 72(a0)
120     ld t1, 80(a0)
121     ld t2, 88(a0)
122     ld s0, 96(a0)
123     ld s1, 104(a0)
124     ld a1, 120(a0)
125     ld a2, 128(a0)
126     ld a3, 136(a0)
127     ld a4, 144(a0)
128     ld a5, 152(a0)
129     ld a6, 160(a0)
130     ld a7, 168(a0)
131     ld s2, 176(a0)
132     ld s3, 184(a0)
133     ld s4, 192(a0)
```

A series of things happenIng

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->kernel/trampoline.S/userret()

```
int uid = getuid();
```

```
89 void
90 usertrapret(void)
91 {
92     struct proc *p = myproc();
93
94     // we're about to switch the destination of traps from
95     // kerneltrap() to usertrap(), so turn off interrupts until
96     // we're back in user space, where usertrap() is correct.
97     intr_off();
98
99     // send syscalls, interrupts, and exceptions to uservec in trampoline.S
100     uint64 trampoline_uservec = TRAMPOLINE + (uservec - trampoline);
101     w_stvec(trampoline_uservec);
102
103     // set up trapframe values that uservec will need when
104     // the process next traps into the kernel.
105     p->trapframe->kernel_satp = r_satp(); // kernel page table
106     p->trapframe->kernel_sp = p->kstack + PGSIZE; // process's kernel stack
107     p->trapframe->kernel_trap = (uint64)usertrap;
108     p->trapframe->kernel_hartid = r_tp(); // hartid for cpuid
109
110     // set up the registers that trampoline.S's sret will use
111     // to get to user space.
112
113     // set S Previous Privilege mode to User.
114     unsigned long x = r_sstatus();
115     x &= ~SSTATUS_SPP; // clear SPP to 0 for user mode
116     x |= SSTATUS_SPIE; // enable interrupts in user mode
117     w_sstatus(x);
118
119     // set S Exception Program Counter to the saved user pc.
120     w_sepc(p->trapframe->epc);
121
122     // tell trampoline.S the user page table to switch to.
123     uint64 satp = MAKE_SATP(p->pagetable);
124
125     // jump to userret in trampoline.S at the top of memory, which
126     // switches to the user page table, restores user registers,
127     // and switches to user mode with sret.
128     uint64 trampoline_userret = TRAMPOLINE + (userret - trampoline);
129     ((void (*)(uint64))trampoline_userret)(satp);
130 }
131
```

kernel/trampoline.S

```
100 .globl userret
101 userret:
102     # userret(pagetable)
103     # called by usertrapret() in trap.c to
104     # switch from kernel to user.
105     # a0: user page table, for satp.
106
107     # switch to the user page table.
108     sfence.vma zero, zero
109     csrw satp, a0
110     sfence.vma zero, zero
111
112     li a0, TRAPFRAME
113
114     # restore all but a0 from TRAPFRAME
115     ld ra, 40(a0)
116     ld sp, 48(a0)
117     ld gp, 56(a0)
118     ld tp, 64(a0)
119     ld t0, 72(a0)
120     ld t1, 80(a0)
121     ld t2, 88(a0)
122     ld s0, 96(a0)
123     ld s1, 104(a0)
124     ld a1, 120(a0)
125     ld a2, 128(a0)
126     ld a3, 136(a0)
127     ld a4, 144(a0)
128     ld a5, 152(a0)
129     ld a6, 160(a0)
130     ld a7, 168(a0)
131     ld s2, 176(a0)
132     ld s3, 184(a0)
133     ld s4, 192(a0)
```

trapframe->a0
contains the
return value

A series of things happenIng

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->kernel/trampoline.S/userret()

```
int uid = getuid();
```

```
89 void
90 usertrapret(void)
91 {
92     struct proc *p = myproc();
93
94     // we're about to switch the destination of traps from
95     // kerneltrap() to usertrap(), so turn off interrupts until
96     // we're back in user space, where usertrap() is correct.
97     intr_off();
98
99     // send syscalls, interrupts, and exceptions to uservec in trampoline.S
100     uint64 trampoline_uservec = TRAMPOLINE + (uservec - trampoline);
101     w_stvec(trampoline_uservec);
102
103     // set up trapframe values that uservec will need when
104     // the process next traps into the kernel.
105     p->trapframe->kernel_satp = r_satp(); // kernel page table
106     p->trapframe->kernel_sp = p->kstack + PGSIZE; // process's kernel stack
107     p->trapframe->kernel_trap = (uint64)usertrap;
108     p->trapframe->kernel_hartid = r_tp(); // hartid for cpuid
109
110     // set up the registers that trampoline.S's sret will use
111     // to get to user space.
112
113     // set S Previous Privilege mode to User.
114     unsigned long x = r_sstatus();
115     x &= ~SSTATUS_SPP; // clear SPP to 0 for user mode
116     x |= SSTATUS_SPIE; // enable interrupts in user mode
117     w_sstatus(x);
118
119     // set S Exception Program Counter to the saved user pc.
120     w_sepc(p->trapframe->epc);
121
122     // tell trampoline.S the user page table to switch to.
123     uint64 satp = MAKE_SATP(p->pagetable);
124
125     // jump to userret in trampoline.S at the top of memory, which
126     // switches to the user page table, restores user registers,
127     // and switches to user mode with sret.
128     uint64 trampoline_userret = TRAMPOLINE + (userret - trampoline);
129     ((void (*)(uint64))trampoline_userret)(satp);
130 }
131
```

kernel/trampoline.S

```
100 .globl userret
101 userret:
102     # userret(pagetable)
103     # called by usertrapret() in trap.c to
104     # switch from kernel to user.
105     # a0: user page table, for satp.
106
107     # switch to the user page table.
108     sfence.vma zero, zero
109     csrw satp, a0
110     sfence.vma zero, zero
111
112     li a0, TRAPFRAME
113
114     # restore all but a0 from TRAPFRAME
115     ld ra, 40(a0)
116     ld sp, 48(a0)
117     ld gp, 56(a0)
118     ld tp, 64(a0)
119     ld t0, 72(a0)
120     ld t1, 80(a0)
121     ld t2, 88(a0)
122     ld <0, 96(a0)
```

trapframe->a0
contains the
return value

```
146 # restore user a0
147 ld a0, 112(a0)
148
149 # return to user mode and user pc.
150 # usertrapret() set up sstatus and sepc.
151 sret
152
153 ld s3, 184(a0)
154 ld s4, 192(a0)
```

A series of things happen [RECAP & NOTE]

user/gibuid.c -> user/Usys.s -> kernel/trampoline.S/uservec() -> kernel/trap.c/usertrap()

```
int uid = getuid(); 108 .global getuid
109 getuid:
110 li a7, SYS_getuid
111 ecall
112 ret
```

```
33 // handle an interrupt, exception, or system call from user space.
34 // called from trampoline.S
35 //
36 void
37 usertrap(void)
38 {
39     int which_dev = 0;
40
41     if((r_sstatus() & SSTATUS_SPP) != 0)
42         panic("usertrap: not from user mode");
43
44     // send interrupts and exceptions to kerneltrap(),
45     // since we're now in the kernel.
46     w_stvec((uint64)kernelvec);
47
48     struct proc *p = myproc();
49
50     // save user program counter.
51     p->trapframe->epc = r_sepc();
52
53     if(r_scause() == 8){
54         // system call
55
56         if(killed(p))
57             exit(-1);
58
59         // sepc points to the ecall instruction,
60         // but we want to return to the next instruction.
61         p->trapframe->epc += 4;
62
63         // an interrupt will change sepc, scause, and sstatus,
64         // so enable only now that we're done with those registers.
65         intr_on();
66
67         syscall();
```

```
19 trampoline:
20 .align 4
21 .globl uservec
22 uservec:
23 #
24 # trap.c sets stvec to point here, so
25 # traps from user space start here,
26 # in supervisor mode, but with a
27 # user page table.
28 #
29
30 # save user a0 in sscratch so
31 # a0 can be used to get at TRAPFRAME.
32 csrw sscratch, a0
33
34 # each process has a separate p->trapframe memory area,
35 # but it's mapped to the same virtual address
36 # (TRAPFRAME) in every process's user page table.
37 li a0, TRAPFRAME
38
39 # save the user registers in TRAPFRAME
40 sd ra, 40(a0)
41 sd sp, 48(a0)
42 sd gp, 56(a0)
43
44 # jump to usertrap(), which does not return
45 jr t0
46
47
48 sd s0, 20(a0)
49 sd s1, 104(a0)
50 sd a1, 120(a0)
51 sd a2, 128(a0)
52 sd a3, 136(a0)
53 sd a4, 144(a0)
54 sd a5, 152(a0)
55 sd a6, 160(a0)
56 sd a7, 168(a0)
57 sd s2, 176(a0)
58 sd s3, 184(a0)
```


A series of things happenIng

user/gibuid.c->user/Usys.s ->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()->kernel/trampoline.S/userret()

```
int uid = getuid();
```

```
89 void
90 usertrapret(void)
91 {
92     struct proc *p = myproc();
93
94     // we're about to switch the destination of traps from
95     // kerneltrap() to usertrap(), so turn off interrupts until
96     // we're back in user space, where usertrap() is correct.
97     intr_off();
98
99     // send syscalls, interrupts, and exceptions to uservec in trampoline.S
100     uint64 trampoline_uservec = TRAMPOLINE + (uservec - trampoline);
101     w_stvec(trampoline_uservec);
102
103     // set up trapframe values that uservec will need when
104     // the process next traps into the kernel.
105     p->trapframe->kernel_satp = r_satp(); // kernel page table
106     p->trapframe->kernel_sp = p->kstack + PGSIZE; // process's kernel stack
107     p->trapframe->kernel_trap = (uint64)usertrap;
108     p->trapframe->kernel_hartid = r_tp(); // hartid for cpuid
109
110     // set up the registers that trampoline.S's sret will use
111     // to get to user space.
112
113     // set S Previous Privilege mode to User.
114     unsigned long x = r_sstatus();
115     x &= ~SSTATUS_SPP; // clear SPP to 0 for user mode
116     x |= SSTATUS_SPIE; // enable interrupts in user mode
117     w_sstatus(x);
118
119     // set S Exception Program Counter to the saved user pc.
120     w_sepc(p->trapframe->epc);
121
122     // tell trampoline.S the user page table to switch to.
123     uint64 satp = MAKE_SATP(p->pagetable);
124
125     // jump to userret in trampoline.S at the top of memory, which
126     // switches to the user page table, restores user registers,
127     // and switches to user mode with sret.
128     uint64 trampoline_userret = TRAMPOLINE + (userret - trampoline);
129     ((void (*)(uint64))trampoline_userret)(satp);
130 }
131
```

kernel/trampoline.S

```
100 .globl userret
101 userret:
102     # userret(pagetable)
103     # called by usertrapret() in trap.c to
104     # switch from kernel to user.
105     # a0: user page table, for satp.
106
107     # switch to the user page table.
108     sfence.vma zero, zero
109     csrw satp, a0
110     sfence.vma zero, zero
111
112     li a0, TRAPFRAME
113
114     # restore all but a0 from TRAPFRAME
115     ld ra, 40(a0)
116     ld sp, 48(a0)
117     ld gp, 56(a0)
118     ld tp, 64(a0)
119     ld t0, 72(a0)
120     ld t1, 80(a0)
121     ld t2, 88(a0)
122     ld s0, 96(a0)
123
```

trapframe->a0
contains the
return value

```
146 # restore user a0
147 ld a0, 112(a0)
148
149 # return to user mode and user pc.
150 # usertrapret() set up sstatus and sepc.
151 sret
152
153 ld s3, 184(a0)
154 ld s4, 192(a0)

```

A series of things happened

user/gibuid.c->user/Usys.s

```
int uid = getuid(); 108  .global getuid
                    109  getuid:
                    110  li a7, SYS_getuid
                    111  ecall
                    112  ret ←
```


A series of things happened (We are back!)

user/gibuid.c

```
#include "kernel/types.h"
#include "kernel/stat.h"
#include "user/user.h"

int main()
{
    int uid = getuid();

    printf("%d\n", uid);

    return 0;
}
```

Okay, but how to pass arguments?

Say, we want to set the uid to 456! (As we know it's simply a variable)

The call will be `setuid(456)`

```
if(!setuid(456)) {  
    printf("setuid success\n");  
} else {  
    printf("setuid failed\n");  
}
```

xv6 stores the arguments passed in registers `a0`, `a1`, `a2`, `a3`, `a4`, `a5`

Yet another system call

user/gibuid.c -> user/Usys.s -> kernel/trampoline.S/uservec() -> kernel/trap.c/usertrap() [*stores registers in trapframe*] -> kernel/syscall.c/syscall()

kernel/syscall.c

```
135 void
136 syscall(void)
137 {
138     int num;
139     struct proc *p = myproc();
140
141     num = p->trapframe->a7;
142     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143         // Use num to lookup the system call function for num, call it,
144         // and store its return value in p->trapframe->a0
145         p->trapframe->a0 = syscalls[num]();
146     } else {
147         printf("%d %s: unknown sys call %d\n",
148             p->pid, p->name, num);
149         p->trapframe->a0 = -1;
150     }
151 }
```

Yet another system call

user/gibuid.c->user/Usys.s->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()*[stores registers in trapframe]*->
kernel/syscall.c/syscall()

kernel/syscall.c

```
static uint64
argraw(int n)
{
    struct proc *p = myproc();
    switch (n) {
        case 0:
            return p->trapframe->a0;
        case 1:
            return p->trapframe->a1;
        case 2:
            return p->trapframe->a2;
        case 3:
            return p->trapframe->a3;
        case 4:
            return p->trapframe->a4;
        case 5:
            return p->trapframe->a5;
    }
    panic("argraw");
    return -1;
}

void
argint(int n, int *ip)
{
    *ip = argraw(n);
}

// Retrieve an argument as a pointer.
// Doesn't check for legality, since
// copyin/copyout will do that.
void
argaddr(int n, uint64 *ip)
{
    *ip = argraw(n);
}

// Fetch the nth word-sized system call
// Copies into buf, at most max.
// Returns string length if OK (including
// null terminator), 0 if buf too small.
int
argstr(int n, char *buf, int max)
{
    uint64 addr;
    argaddr(n, &addr);
    return fetchstr(addr, buf, max);
}
```

kernel/syscall.c

```
135 void
136 syscall(void)
137 {
138     int num;
139     struct proc *p = myproc();
140
141     num = p->trapframe->a7;
142     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
143         // Use num to lookup the system call function for num, call it,
144         // and store its return value in p->trapframe->a0
145         p->trapframe->a0 = syscalls[num]();
146     } else {
147         printf("%d %s: unknown sys call %d\n",
148             p->pid, p->name, num);
149         p->trapframe->a0 = -1;
150     }
151 }
```

Yet another system call

user/gibuid.c->user/Usys.s->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()*[stores registers in trapframe]*->
kernel/syscall.c/syscall() *[say, it called sys_setuid() in sysproc.c]*

kernel/syscall.c

```
static uint64
argraw(int n)
{
    struct proc *p = myproc();
    switch (n) {
    case 0:
        return p->trapframe->a0;
    case 1:
        return p->trapframe->a1;
    case 2:
        return p->trapframe->a2;
    case 3:
        return p->trapframe->a3;
    case 4:
        return p->trapframe->a4;
    case 5:
        return p->trapframe->a5;
    }
    panic("argraw");
    return -1;
}

void
argint(int n, int *ip)
{
    *ip = argraw(n);
}

// Retrieve an argument as a pointer.
// Doesn't check for legality, since
// copyin/copyout will do that.
void
argaddr(int n, uint64 *ip)
{
    *ip = argraw(n);
}

// Fetch the nth word-sized system call argument.
// Copies into buf, at most max.
// Returns string length if OK (including null terminator),
// or -1 if error.
int
argstr(int n, char *buf, int max)
{
    uint64 addr;
    argaddr(n, &addr);
    return fetchstr(addr, buf, max);
}
```

kernel/sysproc.c

```
uint64
sys_setuid(void)
{
    int uid;
    argint(0, &uid);

    if (uid < 0 || uid > 65535) {
        return -1;
    }

    setuid(uid);

    return 0;
}
```

Yet another system call

user/gibuid.c->user/Usys.s->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()*[stores registers in trapframe]*->
kernel/syscall.c/syscall() *[say, it called sys_setuid() in sysproc.c]*

kernel/syscall.c

```
static uint64
argraw(int n)
{
    struct proc *p = myproc();
    switch (n) {
        case 0:
            return p->trapframe->a0;
        case 1:
            return p->trapframe->a1;
        case 2:
            return p->trapframe->a2;
        case 3:
            return p->trapframe->a3;
        case 4:
            return p->trapframe->a4;
        case 5:
            return p->trapframe->a5;
    }
    panic("argraw");
    return -1;
}
```

```
void
argint(int n, int *ip)
{
    *ip = argraw(n);
}

// Retrieve an argument as a pointer.
// Doesn't check for legality, since
// copyin/copyout will do that.
void
argaddr(int n, uint64 *ip)
{
    *ip = argraw(n);
}

// Fetch the nth word-sized system call argument
// Copies into buf, at most max.
// Returns string length if OK (including null)
int
argstr(int n, char *buf, int max)
{
    uint64 addr;
    argaddr(n, &addr);
    return fetchstr(addr, buf, max);
}
```

kernel/sysproc.c

```
uint64
sys_setuid(void)
{
    int uid;
    argint(0, &uid);

    if (uid < 0 || uid > 65535) {
        return -1;
    }

    setuid(uid);

    return 0;
}
```

user/gibuid.c

```
if(!setuid(456)) {
    printf("setuid success\n");
} else {
    printf("setuid failed\n");
}
```

Only **one** argument. For others, n would be 1,2,3,..

Yet another system call

user/gibuid.c->user/Usys.s->kernel/trampoline.S/uservec()->kernel/trap.c/usertrap()*[stores registers in trapframe]*->
kernel/syscall.c/syscall() *[say, it called sys_setuid() in sysproc.c]*

kernel/syscall.c

```
static uint64
argraw(int n)
{
    struct proc *p = myproc();
    switch (n) {
        case 0:
            return p->trapframe->a0;
        case 1:
            return p->trapframe->a1;
        case 2:
            return p->trapframe->a2;
        case 3:
            return p->trapframe->a3;
        case 4:
            return p->trapframe->a4;
        case 5:
            return p->trapframe->a5;
    }
    panic("argraw");
    return -1;
}

void
argint(int n, int *ip)
{
    *ip = argraw(n);
}

// Retrieve an argument as a pointer.
// Doesn't check for legality, since
// copyin/copyout will do that.
void
argaddr(int n, uint64 *ip)
{
    *ip = argraw(n);
}

// Fetch the nth word-sized system call argument.
// Copies into buf, at most max.
// Returns string length if OK (including null terminator),
// or -1 if error.
int
argstr(int n, char *buf, int max)
{
    uint64 addr;
    argaddr(n, &addr);
    return fetchstr(addr, buf, max);
}
```

kernel/sysproc.c

```
uint64
sys_setuid(void)
{
    int uid;
    argint(0, &uid);


    if (uid < 0 || uid > 65535) {
        return -1;
    }

    setuid(uid);

    return 0;
}
```

kernel/proc.c

```
int setuid(int newuid) {
    uid = newuid;
    return 0;
}
```



Thank You



Catalin Pit

@catalinmpit

I'm so glad I did a CS degree
learning Data Structures,
Algorithms, Operating Systems,
Databases, Linear Algebra,
Software Engineering, Information
Analysis, Networking and many
more.

How else would I have been able
to center a div with CSS or change
the background color?